



## Individual Assignment 2

Report and answers of all tasks of second assignment of Opration Systems 1DV512-HT21



*Author:* Saleh Shalabi

*ss225bx@student.lnu.se*

*Semester:* HT21

*Course code:* 1DV512

*Date:* 09/12/2021



1. Does ULE support threads, or does it support processes only? How about CFS?

Both ULE and CFS are done to “schedule large number of threads on large multicore machines”. The answer is yes both of them support threads.

2. How does CFS select the next task to be executed?

The CFS queuing algorithm divides CPU cycles between the threads depending on their priority that is represented by their niceness, the higher niceness the lower priority.

The scheduler in CFS is done by the *vruntime* which can be calculated by dividing the time a thread used from CPU time by its priority. That guarantees that threads with same priority shares core resources fairly. In the scheduler the thread with lowest *vruntime* runs next when a current running thread is preempted.

3. What is a *cgroup* and how is it used by CFS? Does ULE support *cgroups*?

The *cgroup* in CFS are a structure of threads of same application grouped together. That means the *vruntime* of it is the sum of all threads runtime in it. The CFS applies the scheduling algorithm on the *cgroup*, which will ensure fairness between groups.

The ULE does not support the *cgroup* structure.

4. How many queues in total does ULE use? What is the purpose of each queue?

The ULE uses three queues. two to schedule threads. The first one contains interactive threads and the second that contains batch threads. The last one contains only idle task and is used when the core is idle.

Two runqueues is used to give priority to the interactive threads. While the batch is usually used to run without user interaction, which makes its scheduling latency less important.



5. How does ULE compute priority for various tasks?

In the ULE the for the interactive threads the priority is a linear interpolation of their *score*, the less the score is the higher priority it has. In the other hand the batch threads priority is depending on their runtime, the more time a thread spend running the lower priority it has. To get a linear effect on the priority the niceness value is added.

To classify a thread's priority the ULE computes first the *score* of the thread which is defined by *interactivity penalty* + *niceness*. Interactivity penalty is a metric between 0 and 100, and its defined as a function of time *r* a thread spend running and time *s* which is the time a thread spend voluntarily sleeping. If its score is under a certain threshold value then the thread is considered as interactive. The niceness value of 0 is corresponding to spending more than 60% of the time sleeping (not including waiting time for CPU). That function is calculated as following:

$$\text{scaling factor} = m = 50$$

$$\text{penalty}(r,s) = \begin{cases} \frac{m}{\frac{s}{r}} & s > r \\ \frac{m}{\frac{r}{s}} + m & \text{otherwise} \end{cases}$$

Otherwise the thread is classified as a batch. That means negative value of niceness what means higher priority make it easier for a thread to be considered as interactive.

In both interactive and batch runqueues there is a FIFO "first in first out" which is used when a thread is added to the runqueue, the scheduling algorithm insert the added thread at the end of the FIFO and its given an index by the thread priority. Picking a thread to run will be done by taking the first thread in highest priority and non empty FIFO.

In my understanding FIFO is considered as a list in list, all threads with priority 0 will be added in first position (first list in the FIFO list) in the list and so on. For example if we have already 2 threads in it then we add one more thread with priority 0 it will be added in the same position but will be executed when the 2 threads before are executed.



6. Do CFS and ULE support task preemption? Are there any limitations?

The CFS support the task preemption while in ULE is disabled which means that only threads of the kernel can preempt others.

7. Did Bouron et al. discover large differences in per-core scheduling performance between CFS and ULE? Which definition of "performance" did they use in their benchmark, and why?

They found that the main difference between the ULE and CFS is in handling of batch threads. While the CFS want to be fair for all threads the ULE gives the priority to the interactive threads.

For database workloads and NAS they defines the performance of comparing the number of operations per seconds and for other applications they compared as "1/execution time"

They define it in this way because it's that the scheduling algorithm has little influence on the most of workloads, that's because most applications uses threads that perform same work, which means that both CFS and ULE end up scheduling the threads in a round-robin fashion.

8. What is the difference between the multi-core load balancing strategies used by CFS and ULE? Is any of them faster? Does any of them typically reach perfect load balancing?

In both CFS and ULE the workload balancing periodically and placement of threads happens when threads are created or waken up. CFS depend on a complex load matrix that use a hierarchical load strategy to balance. and it runs every 4ms, while the ULE tries to be fair by evening out the number of threads on each core. Balancing in ULE happens less often than CFS and it ignores topology of the hardware.

In Bouron et al. test shows that CFS is way faster in balancing load than the ULE, the ULE takes one thread to each core each balancing workload and while it runs once each (0,5s – 1,5s) less often than the CFS, at the end it reach a perfect balancing state (if the threads are perfectly divisible by the number of cores) while the CFS never reaches a perfect balance because it only tries to balance the load between NUMA nodes when it has an enough big imbalanced (25% load different in practice).



**Source reference:**

All answers is referred to the paper “The Battle of the Schedulers: FreeBSD ULE vs. Linux CFS”

Link: <https://www.usenix.org/conference/atc18/presentation/bouron>